




KEY FEATURES:

- Instance Normalization (batch-size independent)
- Residual Connections (gradient flow)
- Strided Convolution Downsampling (learned)
- Adaptive Masked Transformer (4 layers, 8 heads, 384 dim)
- CBAM Attention on Skip Connections
- Multi-Scale Fusion (all decoder levels)
- ROI-Guided Classification (gradient stopping)
- Deep Supervision (weight 0.3 at 3 scales)

TOTAL PARAMETERS: 37M | INPUT: 256x256x4 | OUTPUTS: Segmentation (3 classes) + Classification (2 classes)